

# Group Assignment Report

## Group Members:

Bruno Simione Beltrame  
Nilton Antonio De Oliveira Paes  
Fabio Cesar Torres De Araujo Hercules  
Oleksii Kachan  
Francisco Isaac De La Rosa Rivero

Dataset of Stock Market in the S&P 500 index which shows the information of all stock available in the index from Feb 2013 to Feb 2018

- Date – day of when the data was taken
- Open – the initial value of the stock
- High – the highest value of the day
- Low – the lowest value of the day
- Close – the closing value of the stock
- Volume – total amount of trades
- Name – stock name

Source of dataset <https://www.kaggle.com/camnugent/sandp500>

Although we can get all the different stocks in the index we decided to predict the closing value of Amazon (AMZN) taking the open value with the max and min of the day.

First we take only the AMZN's stock we create two files one for the training and one for the testing

We load the dataset to train and we take the values to train and to predict in this case is the close column

```
training_data_df =  
pd.read_csv("/dbfs/FileStore/tables/all_stocks_5yr_amazon_train.csv",  
dtype=float)  
X_training = training_data_df.drop('close', axis=1).values  
Y_training = training_data_df[['close']].values
```

And likewise we take the test data and we use the same close column

```
test_data_df =  
pd.read_csv("/dbfs/FileStore/tables/all_stocks_5yr_amazon_test.csv",  
dtype=float)  
X_testing = test_data_df.drop('close', axis=1).values  
Y_testing = test_data_df[['close']].values
```

Taking the range to preprocess the data we take 0 to 1

```
X_scaler = MinMaxScaler(feature_range=(0, 1))  
Y_scaler = MinMaxScaler(feature_range=(0, 1))
```

Using this we going scale the values that was taken previously for both the training set and testing set

```
X_scaled_training = X_scaler.fit_transform(X_training)  
Y_scaled_training = Y_scaler.fit_transform(Y_training)  
  
X_scaled_testing = X_scaler.transform(X_testing)  
Y_scaled_testing = Y_scaler.transform(Y_testing)
```

Now we can define some values to use in the neural network model, first we the define the learning\_rate at which the weight is going to change and how many times we are going to process this using with training\_epochs, in this case we set a step of 0.001 and 200 epochs

```
Learning_rate = 0.001  
training_epochs = 200
```

We currently have 4 inputs and 1 output so we declare as follows

```
number_of_inputs = 4  
number_of_outputs = 1
```

And finally we define the number of nodes per layer

```
Layer_1_nodes = 75  
Layer_2_nodes = 100  
Layer_3_nodes = 75
```

When we have this initial values for the neural network model we now can use them to define the layers

#### Input Layer

```
with tf.variable_scope('input'):
    X = tf.placeholder(tf.float32, shape=(None, number_of_inputs))
```

#### Layer 1

```
with tf.variable_scope('layer_1'):
    weights = tf.get_variable("weights1", shape=[number_of_inputs,
layer_1_nodes], initializer=tf.contrib.layers.xavier_initializer())
    biases = tf.get_variable(name="biases1", shape=[layer_1_nodes],
initializer=tf.zeros_initializer())
    layer_1_output = tf.nn.relu(tf.matmul(X, weights) + biases)
```

#### Layer 2

```
with tf.variable_scope('layer_2'):
    weights = tf.get_variable("weights2", shape=[layer_1_nodes,
layer_2_nodes], initializer=tf.contrib.layers.xavier_initializer())
    biases = tf.get_variable(name="biases2", shape=[layer_2_nodes],
initializer=tf.zeros_initializer())
    layer_2_output = tf.nn.relu(tf.matmul(layer_1_output, weights) + biases)
```

#### Layer 3

```
with tf.variable_scope('layer_3'):
    weights = tf.get_variable("weights3", shape=[layer_2_nodes,
layer_3_nodes], initializer=tf.contrib.layers.xavier_initializer())
    biases = tf.get_variable(name="biases3", shape=[layer_3_nodes],
initializer=tf.zeros_initializer())
    layer_3_output = tf.nn.relu(tf.matmul(layer_2_output, weights) + biases)
```

## Output Layer

```
with tf.variable_scope('output'):
    weights = tf.get_variable("weights4", shape=[layer_3_nodes,
number_of_outputs], initializer=tf.contrib.layers.xavier_initializer())
    biases = tf.get_variable(name="biases4", shape=[number_of_outputs],
initializer=tf.zeros_initializer())
    prediction = tf.matmul(layer_3_output, weights) + biases
```

Later we make the cost that will determine how accurate the prediction is while the training is going

```
with tf.variable_scope('cost'):
    Y = tf.placeholder(tf.float32, shape=(None, 1))
    cost = tf.reduce_mean(tf.squared_difference(prediction, Y))
```

We define the optimizer that will work in the neural network model

```
with tf.variable_scope('train'):
    optimizer = tf.train.AdamOptimizer(learning_rate).minimize(cost)
```

And the summary for the log process on the neural network

```
with tf.variable_scope('logging'):
    tf.summary.scalar('current_cost', cost)
    summary = tf.summary.merge_all()
saver = tf.train.Saver()
```

We all this layer defined we can start the process of the session running TensorFlow

```
with tf.Session() as session:
    session.run(tf.global_variables_initializer())
    training_writer = tf.summary.FileWriter('./Logs/training', session.graph)
    testing_writer = tf.summary.FileWriter('./Logs/testing', session.graph)
    for epoch in range(training_epochs):
        session.run(optimizer, feed_dict={X: X_scaled_training, Y:
Y_scaled_training})
        # Every 5 training steps, Log our progress
        if epoch % 5 == 0:
            training_cost, training_summary = session.run([cost, summary],
feed_dict={X: X_scaled_training, Y:Y_scaled_training})
            testing_cost, testing_summary = session.run([cost, summary],
feed_dict={X: X_scaled_testing, Y:Y_scaled_testing})
            training_writer.add_summary(training_summary, epoch)
            testing_writer.add_summary(testing_summary, epoch)
            print("Epoch: {} - Training Cost: {} Testing Cost:
{}".format(epoch, training_cost, testing_cost))
            final_training_cost = session.run(cost, feed_dict={X:
X_scaled_training, Y: Y_scaled_training})
            final_testing_cost = session.run(cost, feed_dict={X:
X_scaled_testing, Y: Y_scaled_testing})
            print("Final Training cost: {}".format(final_training_cost))
            print("Final Testing cost: {}".format(final_testing_cost))
            Y_predicted_scaled = session.run(prediction, feed_dict={X:
X_scaled_testing})
            Y_predicted = Y_scaler.inverse_transform(Y_predicted_scaled)
            real_earnings = test_data_df['close'].values[0]
            predicted_earnings = Y_predicted[0][0]
            print("The actual close were {}".format(real_earnings))
            print("Our neural network predicted close of
{}".format(predicted_earnings))

    save_path = saver.save(session, "Logs/trained_model.ckpt")
    print("Model saved: {}".format(save_path))
```

And we can see the process during each iterations

```
Epoch: 80 - Training Cost: 8.377773337997496e-05 Testing Cost: 0.0004893926670774817
Epoch: 85 - Training Cost: 8.059653191594407e-05 Testing Cost: 0.0005367605481296778
Epoch: 90 - Training Cost: 7.346444908762351e-05 Testing Cost: 0.000354340358171612
Epoch: 95 - Training Cost: 6.783728167647496e-05 Testing Cost: 0.0004179234674666077
Epoch: 100 - Training Cost: 6.558767199749127e-05 Testing Cost: 0.000383748731110245
```

```
import matplotlib.pyplot as plt

real_earningsAll = test_data_df['close']
fig, ax = plt.subplots()
ax.scatter(real_earningsAll, Y_predicted)
ax.plot([real_earningsAll.min(), real_earningsAll.max()],
        [real_earningsAll.min(),
         real_earningsAll.max()], 'k--', lw=4)
ax.set_xlabel('Measured')
ax.set_ylabel('Predicted')
display(fig)
```


### **Test #1**

Final Training cost: 4.9574733566259965e-05

Final Testing cost: 0.00021704452228732407

The actual close were \$835.77

Our neural network predicted close of \$835.803466796875



Findings:

Using the following parameters, we were able to predict the value **1436.1113** (Actual value: 1416.78)

*Parameters:*

learning\_rate = 0.001  
training\_epochs = 200  
layer\_1\_nodes = 75  
layer\_2\_nodes = 100  
layer\_3\_nodes = 75


**Test #2**

Final Training cost: 1.9963938029832207e-05

Final Testing cost: 7.552812894573435e-05

The actual close were \$835.77

Our neural network predicted close of \$833.1156616210938



*Findings:*

Using the following parameters, we were able to predict the value **1433.7349**

(Actual value: 1416.78)

*Parameters:*


```
learning_rate = 0.001
training_epochs = 800
layer_1_nodes = 1024
layer_2_nodes = 512
layer_3_nodes = 256
```

```
import matplotlib.pyplot as plt
```

```
data_df =
pd.read_csv("/dbfs/FileStore/tables/all_stocks_5yr_amazon_full.csv")
```

```
fig, ax = plt.subplots()
y_plot = data_df['close']
x_plot = data_df['date']
plt.plot(x_plot, y_plot, '-')
```

```
display(fig)
```



Graph containing the evolution of the AMZN stock in 5 years.



Now using MLlib with same dataset

We load the data

```
regressionDataFrame =  
spark.read.csv('/FileStore/tables/all_stocks_5yr_amazon_full.csv',header=True  
, inferSchema = True)
```

```
regressionDataFrame.show(5)
```

date	open	high	low	close	volume	Name
2013-02-08 00:00:00	261.4	265.25	260.555	261.95	3879078	AMZN
2013-02-11 00:00:00	263.2	263.25	256.6	257.21	3403403	AMZN
2013-02-12 00:00:00	259.19	260.16	257.0	258.7	2938660	AMZN
2013-02-13 00:00:00	261.53	269.96	260.3	269.47	5292996	AMZN
2013-02-14 00:00:00	267.37	270.65	265.4	269.24	3462780	AMZN

We get the resilient distributed dataset (RDD) to transform it to labeled point

```
regressionDataRDDDict = regressionDataFrame.rdd
```

```
regressionDataRDDDict = regressionDataFrame.rdd  
regressionDataRDDDict.take(5)
```

```
Out[142]:  
[Row(date=datetime.datetime(2013, 2, 8, 0, 0), open=261.4, high=265.25, low=260.555, close=261.95, volume=3879078,  
Name='AMZN'),  
Row(date=datetime.datetime(2013, 2, 11, 0, 0), open=263.2, high=263.25, low=256.6, close=257.21, volume=3403403,  
Name='AMZN'),  
Row(date=datetime.datetime(2013, 2, 12, 0, 0), open=259.19, high=260.16, low=257.0, close=258.7, volume=2938660,  
Name='AMZN'),  
Row(date=datetime.datetime(2013, 2, 13, 0, 0), open=261.53, high=269.96, low=260.3, close=269.47, volume=5292996,  
Name='AMZN'),  
Row(date=datetime.datetime(2013, 2, 14, 0, 0), open=267.37, high=270.65, low=265.4, close=269.24, volume=3462780,  
Name='AMZN')]
```

Since the RDD has key with values and we only want to have values we use map for that

```
regressionDataRDD = regressionDataFrame.rdd.map(List)
```

```
regressionDataRDD = regressionDataFrame.rdd.map(list)
regressionDataRDD.take(5)
```

```
Out[143]:
[[datetime.datetime(2013, 2, 8, 0, 0),
 261.4,
 265.25,
 260.555,
 261.95,
 3879078,
 'AMZN'],
```

Once we have only the values we can process with LabeledPoint using the close value as the label and the open, high and low as the features.

```
from pyspark.mllib.regression import LabeledPoint
regressionDataLabelPoint = regressionDataRDD.map(lambda data :
LabeledPoint(data[4],data[1:4]))
```

```
regressionDataLabelPoint.take(5)
```

```
Out[145]:
[LabeledPoint(261.95, [261.4,265.25,260.555]),
 LabeledPoint(257.21, [263.2,263.25,256.6]),
 LabeledPoint(258.7, [259.19,260.16,257.0]),
 LabeledPoint(269.47, [261.53,269.96,260.3]),
 LabeledPoint(269.24, [267.37,270.65,265.4])]
```

We going to split this new data set randomly 70% to train and 30% to test

```
regressionLabelPointSplit = regressionDataLabelPoint.randomSplit([0.7,0.3])
```

## Training set

---

```
regressionLabelPointTrainData = regressionLabelPointSplit[0]  
regressionLabelPointTrainData.take(5)
```

Out[147]:

```
[LabeledPoint(257.21, [263.2,263.25,256.6]),  
 LabeledPoint(258.7, [259.19,260.16,257.0]),  
 LabeledPoint(269.47, [261.53,269.96,260.3]),  
 LabeledPoint(269.24, [267.37,270.65,265.4]),  
 LabeledPoint(269.75, [265.91,270.11,264.5])]
```

---

```
regressionLabelPointTrainData.count()
```

Out[148]: 884

---

## Testing set

---

```
regressionLabelPointTestData = regressionLabelPointSplit[1]  
regressionLabelPointTestData.take(5)
```

Out[149]:

```
[LabeledPoint(261.95, [261.4,265.25,260.555]),  
 LabeledPoint(265.09, [267.63,268.92,263.11]),  
 LabeledPoint(266.41, [270.2,274.3,266.371]),  
 LabeledPoint(263.25, [259.4,265.83,256.86]),  
 LabeledPoint(273.11, [265.36,273.3,264.14])]
```

---

```
regressionLabelPointTestData.count()
```

Out[150]: 375

---

Using the Training Data we set it with 200 iterations and a step of 0.00001 to make a model with linear regression to test it later

```
from pyspark.mllib.regression import LinearRegressionWithSGD as LrSGD  
ourModelWithLinearRegression = LrSGD.train(data =  
regressionLabelPointTrainData, iterations = 200, step = 0.00001, intercept =  
True)
```

```
ourModelWithLinearRegression.intercept
```

```
Out[152]: 1.0004267774059894
```

---

```
ourModelWithLinearRegression.weights
```

```
Out[153]: DenseVector([0.3296, 0.3386, 0.3306])
```

---

```
ourModelWithLinearRegression.save(sc, 'dbfs:/FileStore/tables  
/ourModelWithLinearRegressionAmazonStock_201904152024')
```

Now we can test the model and validate how accurate the model is

```
actualDataandLinearRegressionPredictedData =  
regressionLabelPointTestData.map(Lambda data : (float(data.Label) ,  
float(ourModelWithLinearRegression.predict(data.features))))
```

```
actualDataandLinearRegressionPredictedData.take(5)
```

```
Out[160]:  
[(261.95, 263.1058758900237),  
 (265.09, 267.24644204215144),  
 (266.41, 270.9932822821445),  
 (263.25, 261.4215935924259),  
 (273.11, 268.32203675609145)]
```

We calculate the value of root-mean-square error and finally the r2

```
from pyspark.mllib.evaluation import RegressionMetrics as rmtrcs  
ourLinearRegressionModelMetrics = rmtrcs(actualDataandLinearRegressionPredictedData)  
ourLinearRegressionModelMetrics.rootMeanSquaredError
```

```
Out[161]: 5.514359415419885
```

---

```
ourLinearRegressionModelMetrics.r2
```

```
Out[162]: 0.9996232102011257
```

We can validate the model showing that r2 is greater than 0.5 therefore is a good model

## Further data application

### S&P Index Based investment

#### 1. Get daily change

```
1 # calculate daily rate changes
2 def calcPercent(close, open):
3     return ((close - open) / open) * 100
4
5 dataframe = spDataFrame.withColumn('change', calcPercent(spDataFrame['close'], spDataFrame['open']))
6 dataframe.show()
```

▶ (1) Spark Jobs

▶  dataframe: pyspark.sql.dataframe.DataFrame = [date: timestamp, open: double ... 6 more fields]

	date	open	high	low	close	volume	Name	change
	2013-02-08 00:00:00	15.07	15.12	14.63	14.75	8407500	AAL	-2.123424021234242
	2013-02-11 00:00:00	14.89	15.01	14.26	14.46	8882000	AAL	-2.8878441907320327
	2013-02-12 00:00:00	14.45	14.51	14.1	14.27	8126000	AAL	-1.245674740484427
	2013-02-13 00:00:00	14.3	14.94	14.25	14.66	10259500	AAL	2.5174825174825135
	2013-02-14 00:00:00	14.94	14.96	13.16	13.99	31879900	AAL	-6.358768406961174
	2013-02-15 00:00:00	13.93	14.61	13.93	14.5	15628000	AAL	4.091888011486003
	2013-02-19 00:00:00	14.33	14.56	14.08	14.26	11354400	AAL	-0.48848569434752465
	2013-02-20 00:00:00	14.17	14.26	13.15	13.33	14725200	AAL	-5.928016937191248
	2013-02-21 00:00:00	13.62	13.95	12.9	13.37	11922100	AAL	-1.8355359765051396
	2013-02-22 00:00:00	13.57	13.6	13.21	13.57	6071400	AAL	0.0
	2013-02-25 00:00:00	13.6	13.76	13.0	13.02	7186400	AAL	-4.264705882352942
	2013-02-26 00:00:00	13.14	13.42	12.7	13.26	9419000	AAL	0.9132420091324142
	2013-02-27 00:00:00	13.28	13.62	13.18	13.41	7390500	AAL	0.9789156626506084
	2013-02-28 00:00:00	13.49	13.63	13.39	13.43	6143600	AAL	-0.44477390659748334
	2013-03-01 00:00:00	13.37	13.95	13.32	13.61	7376800	AAL	1.7950635751682888
	2013-03-04 00:00:00	13.5	14.07	13.47	13.9	8174800	AAL	2.9629629629629655
	2013-03-05 00:00:00	14.01	14.05	13.71	14.05	7676100	AAL	0.285510349750185
	2013-03-06 00:00:00	14.52	14.68	14.25	14.57	13243200	AAL	0.34435261707989473
	2013-03-07 00:00:00	14.7	14.93	14.5	14.82	9125300	AAL	0.8163265306122517
	2013-03-08 00:00:00	14.99	15.2	14.84	14.92	10593700	AAL	-0.4669779853235509

only showing top 20 rows

#### 2. Reduce columns

```
1 # Select specific columns from dataframe
2 reducedDF = dataframe.select('date', 'Name', 'change')
3 reducedDF.show()
```

▶ (1) Spark Jobs

▶  reducedDF: pyspark.sql.dataframe.DataFrame = [date: timestamp, Name: string ... 1 more fields]

	date	Name	change
	2013-02-08 00:00:00	AAL	-2.123424021234242
	2013-02-11 00:00:00	AAL	-2.8878441907320327
	2013-02-12 00:00:00	AAL	-1.245674740484427
	2013-02-13 00:00:00	AAL	2.5174825174825135
	2013-02-14 00:00:00	AAL	-6.358768406961174
	2013-02-15 00:00:00	AAL	4.091888011486003
	2013-02-19 00:00:00	AAL	-0.48848569434752465
	2013-02-20 00:00:00	AAL	-5.928016937191248
	2013-02-21 00:00:00	AAL	-1.8355359765051396
	2013-02-22 00:00:00	AAL	0.0
	2013-02-25 00:00:00	AAL	-4.264705882352942
	2013-02-26 00:00:00	AAL	0.9132420091324142
	2013-02-27 00:00:00	AAL	0.9789156626506084
	2013-02-28 00:00:00	AAL	-0.44477390659748334
	2013-03-01 00:00:00	AAL	1.7950635751682888
	2013-03-04 00:00:00	AAL	2.9629629629629655
	2013-03-05 00:00:00	AAL	0.285510349750185
	2013-03-06 00:00:00	AAL	0.34435261707989473
	2013-03-07 00:00:00	AAL	0.8163265306122517
	2013-03-08 00:00:00	AAL	-0.4669779853235509

only showing top 20 rows



### 3. Get daily rate change mean

```
1 # get total average
2 from pyspark.sql.functions import mean, col
3
4 groupedDF = reducedDF.groupBy("Name").mean("change").select("Name", "avg(change)")
5 groupedDF = groupedDF.select(col("Name").alias("name"), col("avg(change)").alias("mean"))
6 groupedDF.show()
```

▶ (3) Spark Jobs

▶  groupedDF: pyspark.sql.dataframe.DataFrame = [name: string, mean: double]

```
+----+-----+
| name |          mean |
+----+-----+
| ALXN | -0.03819572050231... |
| GIS  | 0.056748417622980586 |
| K    | 0.030945186822567604 |
| LEN  | -0.02993897166939417 |
| SPGI | 0.09052882241446729 |
| AIV  | 0.05187183560618326 |
| AVY  | 0.04834277090589205 |
| BF.B | 0.04351933136018874 |
| MMM  | 0.07236058899879151 |
| PKI  | 0.07594906776537622 |
| PPG  | 0.0440281761788716 |
| RF   | 0.03740940624892114 |
| AXP  | 0.022083036046032792 |
| CI   | 0.08535447752213957 |
| IRM  | 0.021304890793449105 |
| WEC  | 0.057792724373221784 |
| INFO | 0.13868457521732652 |
| PFG  | 0.05860423437773062 |
| PM   | 0.04034244258475278 |
| SNA  | 0.034608708151195086 |
+----+-----+
```

#### 4. Sort

```
1 # sort data frame
2 sortedDF = groupedDF.sort(col("mean").desc())
3 sortedDF.show()
4
```

##### ► (1) Spark Jobs

►  sortedDF: pyspark.sql.dataframe.DataFrame = [name: string, mean: double]

```
+-----+-----+
| name |          mean |
+-----+-----+
| HPE | 0.1915503478274737 |
| HPQ | 0.1909562023928024 |
| INFO | 0.13868457521732652 |
| BBY | 0.12799839289891546 |
| NVDA | 0.12357063024288786 |
| HUM | 0.11422863658447747 |
| EA | 0.11386175916937684 |
| FISV | 0.10907777744128025 |
| INTU | 0.10733630762579978 |
| RMD | 0.10530298930080509 |
| ANTM | 0.10479424202254987 |
| HII | 0.1031360210136624 |
| TSS | 0.10231258710884215 |
| NTAP | 0.10205940980971352 |
| ABBV | 0.10162815874338921 |
| ACN | 0.09826030744984464 |
| MAR | 0.09806690664753333 |
| VAR | 0.0979728303444835 |
| HCA | 0.09761553318365702 |
| ILMN | 0.09716045291125679 |
+-----+-----+
```

only showing top 20 rows

## 5. Get top N (10 in this example)

```
1 # get top 10
2 top10 = sortedDF.head(10)
3 top10
```

► (1) Spark Jobs

Out[134]:

```
[Row(name='HPE', mean=0.1915503478274737),
 Row(name='HPQ', mean=0.1909562023928024),
 Row(name='INFO', mean=0.13868457521732652),
 Row(name='BBY', mean=0.12799839289891546),
 Row(name='NVDA', mean=0.12357063024288786),
 Row(name='HUM', mean=0.11422863658447747),
 Row(name='EA', mean=0.11386175916937684),
 Row(name='FISV', mean=0.10907777744128025),
 Row(name='INTU', mean=0.10733630762579978),
 Row(name='RMD', mean=0.10530298930080509)]
```

## 6. Calculate index weight

```
1 # calculate total index
2 df_stats = reducedDF.select(
3     mean(col('change')).alias('mean')
4 ).collect()
5 totalMarketWeight = df_stats[0]['mean']
6 totalMarketWeight
```

► (1) Spark Jobs

Out[147]: 0.03527777096687718



Keras dataset


Using this API we can test and train multiple images with different shapes to detect which number it is

```
import tensorflow as tf
(x_train, y_train), (x_test, y_test) = tf.keras.datasets.mnist.Load_data()
```

```
import matplotlib.pyplot as plt
```

You may select anything up to 60,000

```
image_index = 7777
print(y_train[image_index]) # The label is 8
plt.imshow(x_train[image_index], cmap=plt.get_cmap('gray_r'))
display(plt.show())
x_train.shape
```



Reshaping the array to 4-dims so that it can work with the Keras API

```
x_train = x_train.reshape(x_train.shape[0], 28, 28, 1)
x_test = x_test.reshape(x_test.shape[0], 28, 28, 1)
```

```
input_shape = (28, 28, 1)
```

Making sure that the values are float so that we can get decimal points after division

```
x_train = x_train.astype('float32')
```

```
x_test = x_test.astype('float32')
```

Normalizing the RGB codes by dividing it to the max RGB value.

```
x_train /= 255
```

```
x_test /= 255
```

```
print('x_train shape:', x_train.shape)
```

```
print('Number of images in x_train', x_train.shape[0])
```

```
print('Number of images in x_test', x_test.shape[0])
```

```
x_train shape: (60000, 28, 28, 1) Number of images in x_train 60000 Number of  
images in x_test 10000
```

Importing the required Keras modules containing model and layers

```
from keras.models import Sequential
```

```
from keras.layers import Dense, Conv2D, Dropout, Flatten, MaxPooling2D
```

```
# Creating a Sequential Model and adding the Layers
```

```
model = Sequential()
```

```
model.add(Conv2D(28, kernel_size=(3,3), input_shape=input_shape))
```

```
model.add(MaxPooling2D(pool_size=(2, 2)))
```

```
model.add(Flatten()) # Flattening the 2D arrays for fully connected layers
```

```
model.add(Dense(128, activation=tf.nn.relu))
```

```
model.add(Dropout(0.2))
```

```
model.add(Dense(10, activation=tf.nn.softmax))
```

Using TensorFlow backend.

```
model.compile(optimizer='adam',
```

```
              loss='sparse_categorical_crossentropy',
```

```
              metrics=['accuracy'])
```


```
model.fit(x=x_train,y=y_train, epochs=10)
```

```
59072/60000 [=====>.] - ETA: 0s - loss: 0.0203 - acc: 0.9931
59200/60000 [=====>.] - ETA: 0s - loss: 0.0203 - acc: 0.9931
59328/60000 [=====>.] - ETA: 0s - loss: 0.0202 - acc: 0.9931
59456/60000 [=====>.] - ETA: 0s - loss: 0.0202 - acc: 0.9931
59584/60000 [=====>.] - ETA: 0s - loss: 0.0202 - acc: 0.9931
59712/60000 [=====>.] - ETA: 0s - loss: 0.0202 - acc: 0.9931
59840/60000 [=====>.] - ETA: 0s - loss: 0.0202 - acc: 0.9931
59968/60000 [=====>.] - ETA: 0s - loss: 0.0201 - acc: 0.9931
60000/60000 [=====] - 29s 479us/step - loss: 0.0201 - acc: 0.9931
Out[28]: <keras.callbacks.History at 0x7efdf4c86e48>
```

model.evaluate(x\_test, y\_test)

```
7776/10000 [=====>.....] - ETA: 0s
7968/10000 [=====>.....] - ETA: 0s
8256/10000 [=====>.....] - ETA: 0s
8544/10000 [=====>.....] - ETA: 0s
8832/10000 [=====>....] - ETA: 0s
9120/10000 [=====>...] - ETA: 0s
9440/10000 [=====>..] - ETA: 0s
9728/10000 [=====>.] - ETA: 0s
9984/10000 [=====>.] - ETA: 0s
10000/10000 [=====] - 2s 193us/step
Out[29]: [0.06136563551748695, 0.9845]
```

```
image_index = 4444
plt.imshow(x_test[image_index].reshape(28, 28), cmap='Greys')
display(plt.show())
```



```
pred = model.predict(x_test[image_index].reshape(1, 28, 28, 1))  
print(pred.argmax())
```

9

At the end we can predict using the test to evaluate and show which number looks like